

机器学习第三次实验实验报告

实验要求

在这个练习中，需要用 Python 实现 LeNet5 来完成对 MNIST 数据集中 0-9 10 个手写数字的分类。代码只能使用 python 实现，不能使用 PyTorch 或 TensorFlow 框架。

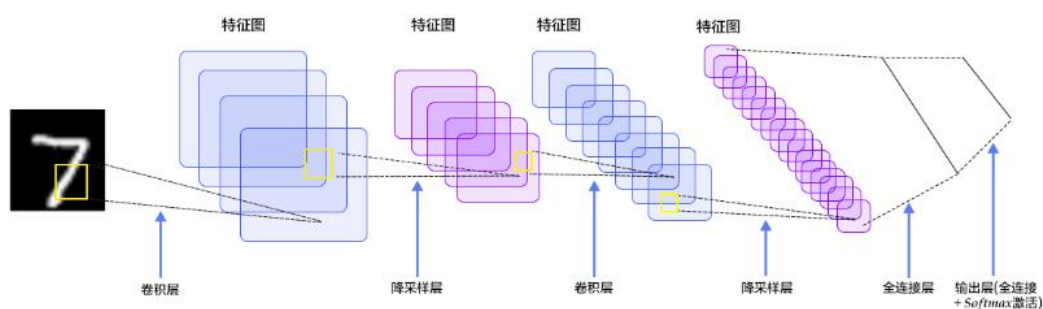
实验环境

Python3.10

Vs Code

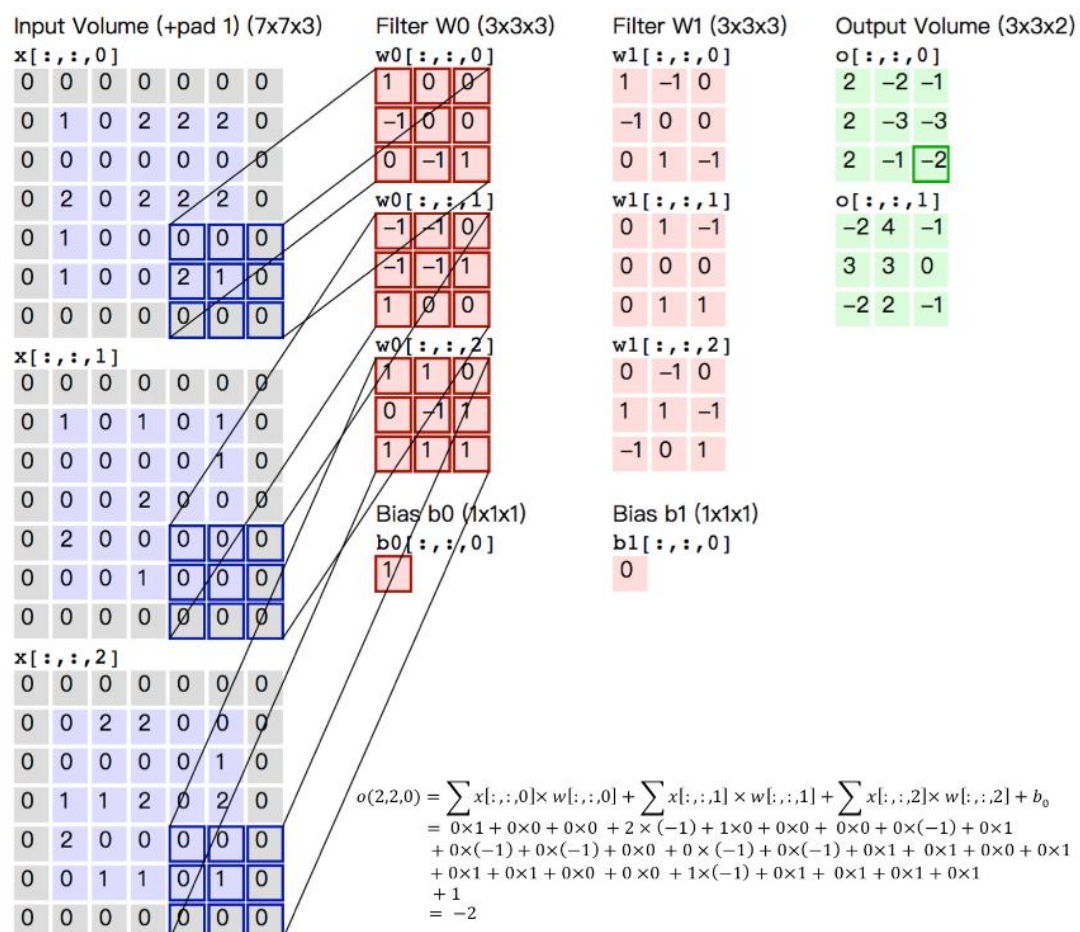
实验准备

LeNet-5 是一个较简单的卷积神经网络。下图显示了其结构：输入的二维图像，先经过两次卷积层到池化层，再经过全连接层，最后使用 softmax 分类作为输出层。



卷积层

卷积层是卷积神经网络的核心基石。在图像识别里我们提到的卷积是二维卷积，即离散二维滤波器（也称作卷积核）与二维图像做卷积操作，简单的讲是二维滤波器滑动到二维图像上所有位置，并在每个位置上与该像素点及其领域像素点做内积。卷积操作被广泛应用与图像处理领域，不同卷积核可以提取不同的特征，例如边沿、线性、角等特征。在深层卷积神经网络中，通过卷积操作可以提取出图像低级到复杂的特征。

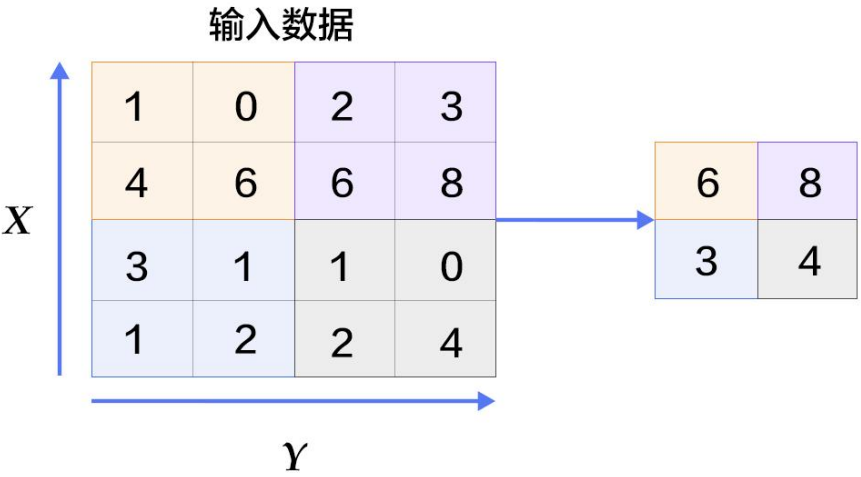


计算过程如上图示例。输入图像大小为 $H=5$ ， $W=5$ ， $D=3$ ，这个示例图中包含两（用 K 表示）组卷积核，即图中滤波器 $W0$ 和 $W1$ 。

在卷积计算中，通常对不同的输入通道采用不同的卷积核，如图示例中每组卷积核包含（ $D=3$ ）个 3×3 （用 $F\times F$ 表示）大小的卷积核。另外，这个示例中卷积核在图像的水平方向（ W 方向）和垂直方向（ H 方向）的滑动步长为 2（用 S 表示）；对输入图像周围各填充 1（用 P 表示）个 0，即图中输入层原始数据为蓝色部分，灰色部分是进行了大小为 1 的扩展，用 0 来进行扩展。经过卷积操作得到输出为 $3\times 3\times 2$ （用 $H_o\times W_o\times K$ 表示）大小的特征图，即 3×3 大小的 2 通道特征图，其中 H_o 计算公式为： $H_o=(H-F+2\times P)/S+1$ ， W_o 同理。而输出特征图中的每个像素，是每组滤波器与输入图像每个特征图的内积再求和，再加上偏置 b_o ，偏置通常对于每个输出特征图是共享的。输出特征图 $o[:, :, 0]$ 中的最后一个 -2 计算如上图右下角公式所示。


在卷积操作中卷积核是可学习的参数，经过上面示例介绍，每层卷积的参数大小为 $D\times F\times F\times K$ 。卷积层的参数较少，这也是由卷积层的主要特性即局部连接和共享权重所决定。

池化层



池化是非线性下采样的一种形式，主要作用是通过减少网络的参数来减小计算量，并且能够在一定程度上控制过拟合。通常在卷积层的后面会加上一个池化层。池化包括最大池化、平均池化等。其中最大池化是用不重叠的矩形框将输入层分成不同的区域，对于每个矩形框的数取最大值作为输出层，如上图所示。

Lenet-5



LeNet-5 共有 7 层，不包含输入，每层都包含可训练参数；每个层有多个 Feature Map，每个 FeatureMap 通过一种卷积滤波器提取输入的一种特征，然后每个 FeatureMap 有多个神经元。

INPUT 层-输入层

首先是数据 INPUT 层，输入图像的尺寸统一归一化为 32*32。

C1 层-卷积层

输入图片：32*32

卷积核大小：5*5

卷积核种类：6

输出 featuremap 大小：28*28 $(32-5+1)=28$

神经元数量: $28 \times 28 \times 6$

可训练参数: $(5 \times 5 + 1) \times 6$ (每个滤波器 $5 \times 5 = 25$ 个 unit 参数和一个 bias 参数, 一共 6 个滤波器)

连接数: $(5 \times 5 + 1) \times 6 \times 28 \times 28 = 122304$

详细说明: 对输入图像进行第一次卷积运算 (使用 6 个大小为 5×5 的卷积核), 得到 6 个 C1 特征图 (6 个大小为 28×28 的 feature maps, $32 - 5 + 1 = 28$)。我们再来看看需要多少个参数, 卷积核的大小为 5×5 , 总共就有 $6 \times (5 \times 5 + 1) = 156$ 个参数, 其中 +1 是表示一个核有一个 bias。对于卷积层 C1, C1 内的每个像素都与输入图像中的 5×5 个像素和 1 个 bias 有连接, 所以总共有 $156 \times 28 \times 28 = 122304$ 个连接(connection)。有 122304 个连接, 但是我们只需要学习 156 个参数, 主要是通过权值共享实现的。

S2 层-池化层 (下采样层)

输入: 28×28

采样区域: 2×2

采样方式: 4 个输入相加, 乘以一个可训练参数, 再加上一个可训练偏置。结果通过 sigmoid

采样种类: 6

输出 featureMap 大小: 14×14 ($28/2$)

神经元数量: $14 \times 14 \times 6$

连接数: $(2 \times 2 + 1) \times 6 \times 14 \times 14$

S2 中每个特征图的大小是 C1 中特征图大小的 $1/4$ 。

详细说明：第一次卷积之后紧接着就是池化运算，使用 2×2 核 进行池化，于是得到了 S2，6 个 14×14 的 特征图 ($28/2=14$)。S2 这个 pooling 层是对 C1 中的 2×2 区域内的像素求和乘以一个权值系数再加上一个偏置，然后将这个结果再做一次映射。同时有 $5 \times 14 \times 14 \times 6 = 5880$ 个连接。

C3 层-卷积层

输入：S2 中所有 6 个或者几个特征 map 组合

卷积核大小： 5×5

卷积核种类：16

输出 featureMap 大小： 10×10 ($14 - 5 + 1 = 10$)

C3 中的每个特征 map 是连接到 S2 中的所有 6 个或者几个特征 map 的，表示本层的特征 map 是上一层提取到的特征 map 的不同组合存在的一个方式是：C3 的前 6 个特征图以 S2 中 3 个相邻的特征图子集为输入。接下来 6 个特征图以 S2 中 4 个相邻特征图子集为输入。然后的 3 个以不相邻的 4 个特征图子集为输入。最后一个将 S2 中所有特征图为输入。

则：可训练参数： $6 \times (3 \times 5 \times 5 + 1) + 6 \times (4 \times 5 \times 5 + 1) + 3 \times (4 \times 5 \times 5 + 1) + 1 \times (6 \times 5 \times 5 + 1)$
 $= 1516$


连接数： $10 \times 10 \times 1516 = 151600$

详细说明：


第一次池化之后是第二次卷积，第二次卷积的输出是 C3，16 个 10x10 的特征图，卷积核大小是 5*5. 我们知道 S2 有 6 个 14*14 的特征图，怎么从 6 个特征图得到 16 个特征图了？这里是通过 S2 的特征图特殊组合计算得到的 16 个特征图。具体如下：

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	X				X	X	X			X	X	X	X		X	X
1	X	X				X	X	X			X	X	X	X		X
2	X	X	X				X	X	X			X		X	X	X
3			X	X	X		X	X	X	X			X		X	X
4				X	X	X		X	X	X	X		X	X		X
5					X	X	X		X	X	X	X		X	X	X

C3 的前 6 个 feature map（对应上图第一个红框的 6 列）与 S2 层相连的 3 个 feature map 相连接（上图第一个红框），后面 6 个 feature map 与 S2 层相连的 4 个 feature map 相连接（上图第二个红框），后面 3 个 feature map 与 S2 层部分不相连的 4 个 feature map 相连接，最后一个与 S2 层的所有 feature map 相连。卷积核大小依然为 5*5，所以总共有 $6 \times (3 \times 5 \times 5 + 1) + 6 \times (4 \times 5 \times 5 + 1) + 3 \times (4 \times 5 \times 5 + 1) + 1 \times (6 \times 5 \times 5 + 1) = 1516$ 个参数。而图像大小为 10*10，所以共有 151600 个连接。



C3 与 S2 中前 3 个图相连的卷积结构如下图所示：



上图对应的参数为 $3 \times 5 \times 5 + 1$ ，一共进行 6 次卷积得到 6 个特征图，所以有 $6 \times (3 \times 5 \times 5 + 1)$ 参数。为什么采用上述这样的组合了？论文中说有两个原因：1) 减少参数，2) 这种不对称的组合连接的方式有利于提取多种组合特征。

S4 层-池化层（下采样层）

输入：10*10

采样区域：2*2

采样方式：4 个输入相加，乘以一个可训练参数，再加上一个可训练偏置。结果通过 sigmoid

采样种类：16

输出 featureMap 大小：5*5 (10/2)

神经元数量：5*5*16=400

连接数：16 * (2*2+1) * 5*5=2000

S4 中每个特征图的大小是 C3 中特征图大小的 1/4

详细说明: S4 是 pooling 层, 窗口大小仍然是 2×2 , 共计 16 个 feature map, C3 层的 16 个 10×10 的图分别进行以 2×2 为单位的池化得到 16 个 5×5 的特征图。有 $5 \times 5 \times 5 \times 16 = 2000$ 个连接。连接的方式与 S2 层类似。

C5 层-卷积层

输入: S4 层的全部 16 个单元特征 map (与 s4 全相连)


卷积核大小: 5×5

卷积核种类: 120

输出 featureMap 大小: 1×1 ($5 - 5 + 1$)

可训练参数/连接: $120 * (16 * 5 * 5 + 1) = 48120$

详细说明: C5 层是一个卷积层。由于 S4 层的 16 个图的大小为 5×5 , 与卷积核的大小相同, 所以卷积后形成的图的大小为 1×1 。这里形成 120 个卷积结果。每个都与上一层的 16 个图相连。所以共有 $(5 \times 5 \times 16 + 1) \times 120 = 48120$ 个参数, 同样有 48120 个连接。C5 层的网络结构如下:



F6 层-全连接层

输入: c5 120 维向量


计算方式: 计算输入向量和权重向量之间的点积, 再加上一个偏置, 结果通过 sigmoid 函数输出。

可训练参数: $84 \times (120 + 1) = 10164$

详细说明: 6 层是全连接层。F6 层有 84 个节点, 对应于一个 7x12 的比特图, -1 表示白色, 1 表示黑色, 这样每个符号的比特图的黑白色就对应于一个编码。该层的训练参数和连接数是 $(120 + 1) \times 84 = 10164$ 。ASCII 编码图如下:



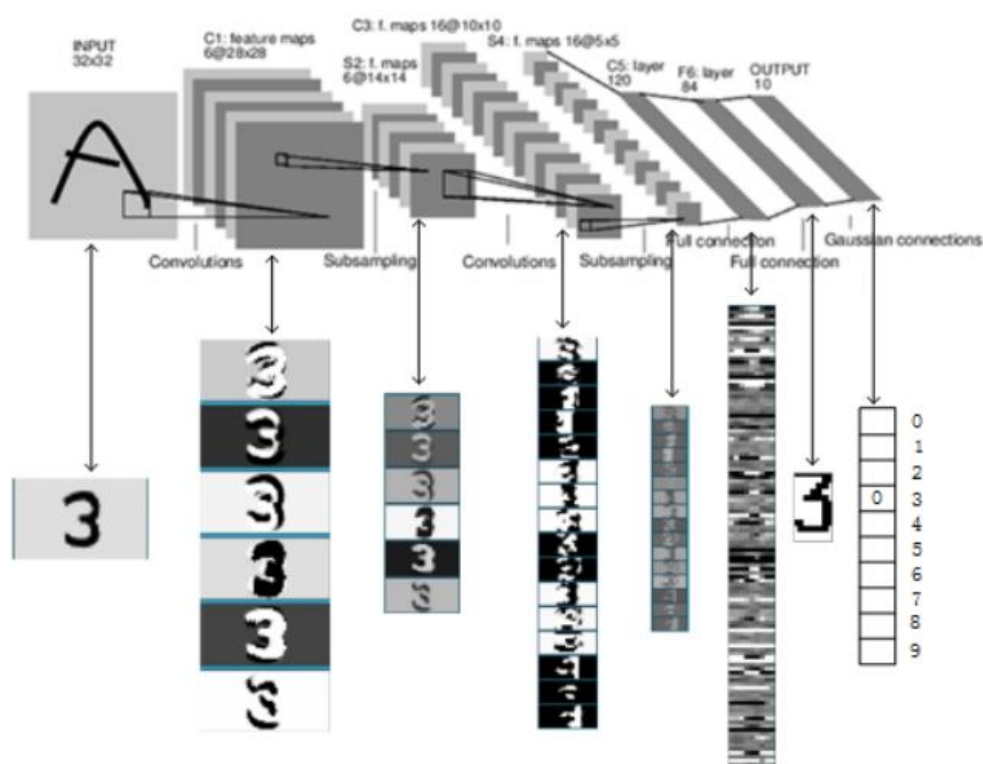
F6 层的连接方式如下:



Output 层-全连接层

Output 层也是全连接层，共有 10 个节点，分别代表数字 0 到 9，且如果节点 i 的值为 0，则网络识别的结果是数字 i 。采用的是径向基函数（RBF）的网络连接方式。假设 x 是上一层的输入， y 是 RBF 的输出，则 RBF 输出的计算方式是：
$$y_i = \sum_j (x_j - w_{ij})^2$$

上式 w_{ij} 的值由 i 的比特图编码确定， i 从 0 到 9， j 取值从 0 到 $7 \times 12 - 1$ 。RBF 输出的值越接近于 0，则越接近于 i ，即越接近于 i 的 ASCII 编码图，表示当前网络输入的识别结果是字符 i 。该层有 $84 \times 10 = 840$ 个参数和连接。



实验内容

Main

main 函数中，首先是对数据的读取和正则化，由于 mnist 数据集中的图片为 $28 * 28$ ，需要 padding 到 $32 * 32$ 。

然后，将处理之后的数据输入 train 函数，开始训练。

训练完成后，在测试集上测试训练网络的性能。

```
main.py > ...
1  # coding=utf-8
2  from data_processing import *
3  from train import *
4  from LeNet5 import *
5
6  if __name__ == '__main__':
7      # get the data
8      train_images, train_labels, test_images, test_labels = load_data()
9
10     print("Got data...\n")
11
12     # data processing, normalization&zero-padding
13     print("Normalization and zero-padding...\n")
14     train_images = normalize(zero_pad(train_images[:, :, :, np.newaxis], 2), 'LeNet5')
15     test_images = normalize(zero_pad(test_images[:, :, :, np.newaxis], 2), 'LeNet5')
16     print("The shape of training image with padding: ", train_images.shape)
17     print("The shape of testing image with padding: ", test_images.shape)
18     print("Finish data processing...\n")
19
20     # train LeNet5
21     LeNet5 = LeNet5()
22     print("Start training...")
23     start_time = time.time()
24     train(LeNet5, train_images, train_labels)
25     end_time = time.time()
26     print("Finished training, the total training time is {}s \n".format(end_time - start_time))
27
28     # read model
29     # with open('model_data_13.pkl', 'rb') as input_:
30     #     LeNet5 = pickle.load(input_)
31
32     # evaluate on test dataset
33     print("Start testing...")
34     error01, class_pred = LeNet5.Forward_Propagation(test_images, test_labels, 'test')
35     print("error rate:", error01 / len(class_pred))
36     print("Finished testing, the accuracy is {} \n".format(1 - error01 / len(class_pred)))
37
```

数据处理 (data_processing)

```
data_processing.py >...
1  # coding=utf-8
2  import numpy as np
3  import struct
4  import os
5
6  data_dir = "D:/wtx/machine-learning/ex1/mnist_data/"
7  train_data_dir = "train-images-idx3-ubyte"
8  train_label_dir = "train-labels-idx1-ubyte"
9  test_data_dir = "t10k-images-idx3-ubyte"
10 test_label_dir = "t10k-labels-idx1-ubyte"
11
12
13 # Load the MNIST data for this exercise
14 def load_mnist(file_dir, is_images=True):
15     # Read binary data
16     bin_file = open(file_dir, 'rb')
17     bin_data = bin_file.read()
18     bin_file.close()
19     # Analysis file header
20     if is_images:
21         # Read images
22         fmt_header = '>iiii'
23         magic, num_images, num_rows, num_cols = struct.unpack_from(fmt_header, bin_data, 0)
24         data_size = num_images * num_rows * num_cols
25         mat_data = struct.unpack_from('>' + str(data_size) + 'B', bin_data, struct.calcsize(fmt_header))
26         mat_data = np.reshape(mat_data, [num_images, num_rows, num_cols])
27     else:
28         # Read labels
29         fmt_header = '>ii'
30         magic, num_images = struct.unpack_from(fmt_header, bin_data, 0)
31         mat_data = struct.unpack_from('>' + str(num_images) + 'B', bin_data, struct.calcsize(fmt_header))
32         mat_data = np.reshape(mat_data, [num_images])
33     print('Load images from %s, number: %d, data shape: %s' % (file_dir, num_images, str(mat_data.shape)))
34     return mat_data
```

```
37 # call the load_mnist function to get the images and labels of training set and testing set
38 def load_data():
39     print('Loading MNIST data from files...')
40     train_images = load_mnist(os.path.join(data_dir, train_data_dir), True)
41     train_labels = load_mnist(os.path.join(data_dir, train_label_dir), False)
42     test_images = load_mnist(os.path.join(data_dir, test_data_dir), True)
43     test_labels = load_mnist(os.path.join(data_dir, test_label_dir), False)
44     return train_images, train_labels, test_images, test_labels
45
46
47 # transfer the image from gray to binary and get the one-hot style labels
48 def data_convert(x, y, m, k):
49     x[x <= 40] = 0
50     x[x > 40] = 1
51     ont_hot_y = np.zeros((m, k))
52     for t in range(m):
53         ont_hot_y[t, y[t]] = 1
54     return x, ont_hot_y
55
```

```
57 # padding for the matrix of images
58 def zero_pad(X, pad):
59     X_pad = np.pad(X, ((0, 0), (pad, pad), (pad, pad), (0, 0)), 'constant', constant_values=(0, 0))
60     return X_pad
61
62
63 # normalization of the input images
64 def normalize(image, mode='LeNet5'):
65     image -= image.min()
66     image = image / image.max()
67     if mode == '0p1':
68         return image # range = [0,1]
69     elif mode == 'n1p1':
70         image = image * 2 - 1 # range = [-1,1]
71     elif mode == 'LeNet5':
72         image = image * 1.275 - 0.1 # range = [-0.1,1.175]
73     return image
74
```


LeNet5

LeNet5 类的 Forward_Propagation 和 Back_Propagation 函数会调用每一层对应的 forward_propagation 和 back_propagation 函数,可以看到,层与层之间的参数会一层层的进行传递。通过 SDLM 算法进行计算得到的每一轮的学习率。

```
LeNet5.py > LeNet5 > _init_
1  import numpy as np
2
3  from pooling_utils import *
4  from convolution_utils import *
5  from activation_utils import *
6  from RBF_init import *
7
8  bitmap = rbf_init_weight()
9
10
11 class LeNet5(object):
12     """
13     C1 -> S2 -> C3 -> S4 -> C5 -> F6 -> Output
14
15     Reference: https://www.cnblogs.com/fengff/p/10173071.html
16     """
17
18     def __init__(self):
19         # Designate combination of kernels and feature maps of S2.
20         C3_mapping = [[0, 1, 2], [1, 2, 3], [2, 3, 4], [3, 4, 5], [4, 5, 0], [5, 0, 1],
21                       [0, 1, 2, 3], [1, 2, 3, 4], [2, 3, 4, 5], [3, 4, 5, 0], [4, 5, 0, 1],
22                       [5, 0, 1, 2], [0, 1, 3, 4], [1, 2, 4, 5], [0, 2, 3, 5], [0, 1, 2, 3, 4, 5]]
23
24         kernel_shape = {"C1": (5, 5, 1, 6),
25                           "C3": (5, 5, 6, 16),
26                           "C5": (5, 5, 16, 120),
27                           "F6": (120, 84),
28                           "OUTPUT": (84, 10)}
29
30         hyper_parameters_convolution = {"stride": 1, "pad": 0}
31         hyper_parameters_pooling = {"stride": 2, "f": 2}
32
33         self.C1 = ConvolutionLayer(kernel_shape["C1"], hyper_parameters_convolution)
34         self.a1 = Activation("LeNet5_squash")
35         self.S2 = PoolingLayer(hyper_parameters_pooling, "average")
36
37         self.C3 = ConvolutionLayer_maps(kernel_shape["C3"], hyper_parameters_convolution, C3_mapping)
38         self.a2 = Activation("LeNet5_squash")
39         self.S4 = PoolingLayer(hyper_parameters_pooling, "average")
40
41         self.C5 = ConvolutionLayer(kernel_shape["C5"], hyper_parameters_convolution)
42         self.a3 = Activation("LeNet5_squash")
43
44         self.F6 = FCLayer(kernel_shape["F6"])
45         self.a4 = Activation("LeNet5_squash")
46
47         self.Output = RBFLayer(bitmap)
48
```

```

49 def Forward_Propagation(self, input_image, input_label, mode):
50     self.label = input_label
51     self.C1_FP = self.C1.forward_propagation(input_image)
52     self.a1_FP = self.a1.forward_propagation(self.C1_FP)
53     self.S2_FP = self.S2.forward_propagation(self.a1_FP)
54
55     self.C3_FP = self.C3.forward_propagation(self.S2_FP)
56     self.a2_FP = self.a2.forward_propagation(self.C3_FP)
57     self.S4_FP = self.S4.forward_propagation(self.a2_FP)
58
59     self.C5_FP = self.C5.forward_propagation(self.S4_FP)
60     self.a3_FP = self.a3.forward_propagation(self.C5_FP)
61
62     self.flatten = self.a3_FP[:, 0, 0, :]
63     self.F6_FP = self.F6.forward_propagation(self.flatten)
64     self.a4_FP = self.a4.forward_propagation(self.F6_FP)
65
66     # output sum of the loss over mini-batch when mode = 'train'
67     # output tuple of (0/1 error, class_predict) when mode = 'test'
68     out = self.Output.forward_propagation(self.a4_FP, input_label, mode)
69
70     return out

```

```

72 def Back_Propagation(self, momentum, weight_decay):
73     dy_pred = self.Output.back_propagation()
74
75     dy_pred = self.a4.back_propagation(dy_pred)
76     F6_BP = self.F6.back_propagation(dy_pred, momentum, weight_decay)
77     reverse_flatten = F6_BP[:, np.newaxis, np.newaxis, :]
78
79     reverse_flatten = self.a3.back_propagation(reverse_flatten)
80     C5_BP = self.C5.back_propagation(reverse_flatten, momentum, weight_decay)
81
82     S4_BP = self.S4.back_propagation(C5_BP)
83     S4_BP = self.a2.back_propagation(S4_BP)
84     C3_BP = self.C3.back_propagation(S4_BP, momentum, weight_decay)
85
86     S2_BP = self.S2.back_propagation(C3_BP)
87     S2_BP = self.a1.back_propagation(S2_BP)
88     C1_BP = self.C1.back_propagation(S2_BP, momentum, weight_decay)
89
90     # Stochastic Diagonal Levenberg-Marquardt method for determining the learning rate before the beginning of each ep
91     def SDLM(self, mu, lr_global):
92         d2y_pred = self.Output.SDLM()
93         d2y_pred = self.a4.SDLM(d2y_pred)]
94
95         F6_SDLM = self.F6.SDLM(d2y_pred, mu, lr_global)
96         reverse_flatten = F6_SDLM[:, np.newaxis, np.newaxis, :]
97
98         reverse_flatten = self.a3.SDLM(reverse_flatten)
99         C5_SDLM = self.C5.SDLM(reverse_flatten, mu, lr_global)
100
101         S4_SDLM = self.S4.SDLM(C5_SDLM)
102         S4_SDLM = self.a2.SDLM(S4_SDLM)
103         C3_SDLM = self.C3.SDLM(S4_SDLM, mu, lr_global)
104
105         S2_SDLM = self.S2.SDLM(C3_SDLM)
106         S2_SDLM = self.a1.SDLM(S2_SDLM)
107         C1_SDLM = self.C1.SDLM(S2_SDLM, mu, lr_global)
108

```

卷积层前向传播和后向传播 (convolution_utils)

```
5 # Numpy version: compute with np.tensordot()
6 def conv_forward(A_prev, W, b, hyper_parameters):
7     """
8     Implements the forward propagation for a convolution function
9
10    :param A_prev: output activations of the previous layer, numpy array of shape (m, n_H_prev, n_W_prev, n_C_prev)
11    :param W: Weights, numpy array of shape (f, f, n_C_prev, n_C)
12    :param b: Biases, numpy array of shape (1, 1, 1, n_C)
13    :param hyper_parameters: python dictionary containing "stride" and "pad"
14
15    :return: Z -- conv output, numpy array of shape (m, n_H, n_W, n_C)
16            cache -- cache of values needed for the conv_backward() function
17    """
18    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
19    (f, f, n_C_prev, n_C) = W.shape
20
21    stride = hyper_parameters["stride"]
22    pad = hyper_parameters["pad"]
23
24    n_H = int((n_H_prev + 2 * pad - f) / stride + 1)
25    n_W = int((n_W_prev + 2 * pad - f) / stride + 1)
26
27    # Initialize the output volume Z with zeros.
28    Z = np.zeros((m, n_H, n_W, n_C))
29    A_prev_pad = zero_pad(A_prev, pad)
30    for h in range(n_H): # loop over vertical axis of the output volume
31        for w in range(n_W): # loop over horizontal axis of the output volume
32            # Use the corners to define the (3D) slice of a_prev_pad.
33            A_slice_prev = A_prev_pad[:, h * stride:h * stride + f, w * stride:w * stride + f, :]
34            # Convolve the (3D) slice with the correct filter W and bias b, to get back one output neuron.
35            Z[:, h, w, :] = np.tensordot(A_slice_prev, W, axes=([1, 2, 3], [0, 1, 2])) + b
36
37    assert (Z.shape == (m, n_H, n_W, n_C))
38    cache = (A_prev, W, b, hyper_parameters)
39    return Z, cache
40
```

```
42 # Numpy version: compute with np.dot
43 def conv_backward(dZ, cache):
44     """
45     Implement the backward propagation for a convolution function
46
47     :param dZ: gradient of the cost with respect to the output of the conv layer (Z), numpy array of shape (m, n_H, n_W, n_C)
48     :param cache: cache of values needed for the conv_backward(), output of conv_forward()
49
50     :return: dA_prev -- gradient of the cost with respect to the input of the conv layer (A_prev),
51             numpy array of shape (m, n_H_prev, n_W_prev, n_C_prev)
52             dW -- gradient of the cost with respect to the weights of the conv layer (W)
53             numpy array of shape (f, f, n_C_prev, n_C)
54             db -- gradient of the cost with respect to the biases of the conv layer (b)
55             numpy array of shape (1, 1, 1, n_C)
56     """
57     (A_prev, W, b, hyper_parameters) = cache
58     (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
59     (f, f, n_C_prev, n_C) = W.shape
60     (m, n_H, n_W, n_C) = dZ.shape
61     stride = hyper_parameters["stride"]
62     pad = hyper_parameters["pad"]
63
64     dA_prev = np.zeros((m, n_H_prev, n_W_prev, n_C_prev))
65     dW = np.zeros((f, f, n_C_prev, n_C))
66     db = np.zeros((1, 1, 1, n_C))
67
68     if pad != 0:
69         A_prev_pad = zero_pad(A_prev, pad)
70         dA_prev_pad = zero_pad(dA_prev, pad)
71     else:
72         A_prev_pad = A_prev
73         dA_prev_pad = dA_prev
```



```

75     for h in range(n_H): # loop over vertical axis of the output volume
76         for w in range(n_W): # loop over horizontal axis of the output volume
77             # Find the corners of the current "slice"
78             vert_start, horiz_start = h * stride, w * stride
79             vert_end, horiz_end = vert_start + f, horiz_start + f
80
81             # Use the corners to define the slice from a_prev_pad
82             A_slice = A_prev_pad[:, vert_start:vert_end, horiz_start:horiz_end, :]
83
84             # Update gradients for the window and the filter's parameters
85             dA_prev_pad[:, vert_start:vert_end, horiz_start:horiz_end, :] += np.transpose(np.dot(W, dZ[:, h, w, :].T), (3, 0, 1, 2))
86
87             dW += np.dot(np.transpose(A_slice, (1, 2, 3, 0)), dZ[:, h, w, :])
88             db += np.sum(dZ[:, h, w, :], axis=0)
89
90     # Set dA_prev to the unpadded dA_prev_pad
91     dA_prev = dA_prev_pad if pad == 0 else dA_prev_pad[:, pad:-pad, pad:-pad, :]
92
93     # Making sure your output shape is correct
94     assert (dA_prev.shape == (m, n_H_prev, n_W_prev, n_C_prev))
95
96     return dA_prev, dW, db
97

```

```

99     def conv_SDLM(dZ, cache):
100         (A_prev, W, b, hparameters) = cache
101         (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
102         (f, n_C_prev, n_W, n_H) = W.shape
103         stride = hparameters["stride"]
104         pad = hparameters["pad"]
105         (m, n_H, n_W, n_C) = dZ.shape
106         dA_prev = np.zeros((m, n_H_prev, n_W_prev, n_C_prev))
107         dW = np.zeros((f, n_C_prev, n_C))
108         db = np.zeros((1, 1, 1, n_C))
109
110         if pad != 0:
111             A_prev_pad = zero_pad(A_prev, pad)
112             dA_prev_pad = zero_pad(dA_prev, pad)
113         else:
114             A_prev_pad = A_prev
115             dA_prev_pad = dA_prev
116
117         for h in range(n_H): # loop over vertical axis of the output volume
118             for w in range(n_W): # loop over horizontal axis of the output volume
119                 # Find the corners of the current "slice"
120                 vert_start, horiz_start = h * stride, w * stride
121                 vert_end, horiz_end = vert_start + f, horiz_start + f
122
123                 # Use the corners to define the slice from a_prev_pad
124                 A_slice = A_prev_pad[:, vert_start:vert_end, horiz_start:horiz_end, :]
125
126                 # Update gradients for the window and the filter's parameters
127                 dA_prev_pad[:, vert_start:vert_end, horiz_start:horiz_end, :] += np.transpose(
128                     np.dot(np.power(W, 2), dZ[:, h, w, :].T), (3, 0, 1, 2))
129
130                 dW += np.dot(np.transpose(np.power(A_slice, 2), (1, 2, 3, 0)), dZ[:, h, w, :])
131
132         # Set dA_prev to the unpadded dA_prev_pad
133         dA_prev = dA_prev_pad if pad == 0 else dA_prev_pad[:, pad:-pad, pad:-pad, :]
134         assert (dA_prev.shape == (m, n_H_prev, n_W_prev, n_C_prev))
135         return dA_prev, dW

```

池化层前向传播和后向传播 (pooling_utils)

```
4 def pool_forward(A_prev, hyper_parameters, mode):
5     m, n_H_prev, n_W_prev, n_C_prev = A_prev.shape
6     f = hyper_parameters["f"]
7     stride = hyper_parameters["stride"]
8
9     n_H = int(1 + (n_H_prev - f) / stride)
10    n_W = int(1 + (n_W_prev - f) / stride)
11    n_C = n_C_prev
12
13    A = np.zeros((m, n_H, n_W, n_C))
14    for h in range(n_H): # loop on the vertical axis of the output volume
15        for w in range(n_W): # loop on the horizontal axis of the output volume
16            # Use the corners to define the current slice on the ith training example of A_prev, channel c
17            A_prev_slice = A_prev[:, h * stride:h * stride + f, w * stride:w * stride + f, :]
18            # Compute the pooling operation on the slice. Use an if statement to differentiate the modes.
19            if mode == "max":
20                A[:, h, w, :] = np.max(A_prev_slice, axis=(1, 2))
21            elif mode == "average":
22                A[:, h, w, :] = np.average(A_prev_slice, axis=(1, 2))
23
24    cache = (A_prev, hyper_parameters)
25    assert (A.shape == (m, n_H, n_W, n_C))
26    return A, cache

```

```
29 def pool_backward(dA, cache, mode):
30     A_prev, hyper_parameters = cache
31
32     stride = hyper_parameters["stride"]
33     f = hyper_parameters["f"]
34
35     m, n_H_prev, n_W_prev, n_C_prev = A_prev.shape # 256,28,28,6
36     m, n_H, n_W, n_C = dA.shape # 256,14,14,6
37
38     dA_prev = np.zeros((m, n_H_prev, n_W_prev, n_C_prev)) # 256,28,28,6
39
40
41     for h in range(n_H): # loop on the vertical axis
42         for w in range(n_W): # loop on the horizontal axis
43             # Find the corners of the current "slice"
44             vert_start, horiz_start = h * stride, w * stride
45             vert_end, horiz_end = vert_start + f, horiz_start + f
46
47             # Compute the backward propagation in both modes.
48             if mode == "max":
49                 A_prev_slice = A_prev[:, vert_start:vert_end, horiz_start:horiz_end, :]
50                 A_prev_slice = np.transpose(A_prev_slice, (1, 2, 3, 0))
51                 mask = A_prev_slice == A_prev_slice.max((0, 1))
52                 mask = np.transpose(mask, (3, 2, 0, 1))
53                 dA_prev[:, vert_start:vert_end, horiz_start:horiz_end, :] \
54                     += np.transpose(np.multiply(dA[:, h, w, :][:, :, np.newaxis, np.newaxis], mask), (0, 2, 3, 1))
55             elif mode == "average":
56                 da = dA[:, h, w, :][:, np.newaxis, np.newaxis, :] # 256*1*1*6
57                 dA_prev[:, vert_start:vert_end, horiz_start:horiz_end, :] += np.repeat(np.repeat(da, 2, axis=1), 2, axis=2) / f / f
58
59     assert (dA_prev.shape == A_prev.shape)
60     return dA_prev

```

```
71 def subsampling_forward(A_prev, weight, b, hparameters):
72     A_, cache = pool_forward(A_prev, hparameters, 'average')
73     A = A_ * weight + b
74     cache_A = (cache, A_)
75     return A, cache_A
76
77
78 def subsampling_backward(dA, weight, b, cache_A):
79     (cache, A_) = cache_A
80     db = dA
81     dW = np.sum(np.multiply(dA, A_))
82     dA_ = dA * weight
83     dA_prev = pool_backward(dA_, cache, 'average')
84     return dA_prev, dW, db

```

训练 (train)

```
train.py ? ...
1  # coding=utf-8
2  import sys
3  import time
4  import math
5  import pickle
6  import numpy as np
7  import matplotlib.pyplot as plt
8
9
10 # Number of epochs & learning rate in the original paper
11 epochs_original, lr_global_original = 16, np.array([5e-4] * 2 + [2e-4] * 3 + [1e-4] * 3 + [5e-5] * 4 + [1e-5] * 8)
12 # Number of epochs & learning rate I used
13 epochs, lr_global_list = epochs_original, lr_global_original * 100
14
```

```
16 def train(LeNet5, train_images, train_labels):
17     momentum = 0.9
18     weight_decay = 0
19     batch_size = 256
20
21     # Training loops
22     cost_last, count = np.Inf, 0
23     err_rate_list = []
24     for epoch in range(0, epochs):
25         print("----- epoch{} begin -----".format(epoch + 1))
26
27         # Stochastic Diagonal Levenberg-Marquardt method for determining the learning rate
28         batch_image, batch_label = random_mini_batches(train_images, train_labels, mini_batch_size=500, one_batch=True)
29         LeNet5.Forward_Propagation(batch_image, batch_label, 'train')
30         lr_global = lr_global_list[epoch]
31         LeNet5.SDLM(0.02, lr_global)
32
33         # print info
34         print("global learning rate:", lr_global)
35         print("learning rates in trainable layers:", np.array([LeNet5.C1.lr, LeNet5.C3.lr, LeNet5.C5.lr, LeNet5.F6.lr]))
36         print("batch size:", batch_size)
37         print("momentum:", momentum, ", weight decay:", weight_decay)
```

```
40     ste = time.time()
41     cost = 0
42     mini_batches = random_mini_batches(train_images, train_labels, batch_size)
43     for i in range(len(mini_batches)):
44         batch_image, batch_label = mini_batches[i]
45
46         loss = LeNet5.Forward_Propagation(batch_image, batch_label, 'train')
47         cost += loss
48
49         LeNet5.Back_Propagation(momentum, weight_decay)
50
51         # print progress
52         if i % (int(len(mini_batches) / 100)) == 0:
53             sys.stdout.write("\033[F") # CURSOR_UP_ONE
54             sys.stdout.write("\033[K") # ERASE_LINE
55             print("progress:", int(100 * (i + 1) / len(mini_batches)), "%", "cost =", cost, end='\r')
56     sys.stdout.write("\033[F") # CURSOR_UP_ONE
57     sys.stdout.write("\033[K") # ERASE_LINE
58
59     print("Done, cost of epoch", epoch + 1, ":", cost, " ")
60
61     error01_train, _ = LeNet5.Forward_Propagation(train_images, train_labels, 'test')
62     err_rate_list.append(error01_train / 60000)
63     # error01_test, _ = LeNet5.Forward_Propagation(test_images, test_labels, 'test')
64     # err_rate_list.append([error01_train / 60000, error01_test / 10000])
65     print("0/1 error of training set:", error01_train, "/", len(train_labels))
66     # print("0/1 error of testing set: ", error01_test, "/", len(test_labels))
67     print("Time used: ", time.time() - ste, "sec")
68     print(
69         "----- epoch{} end ----- \n".format(epoch + 1))
70
71     # conserve the model
72     # with open('model_data_' + str(epoch) + '.pkl', 'wb') as output:
73     #     pickle.dump(LeNet5, output, pickle.HIGHEST_PROTOCOL)
74
75     err_rate_list = np.array(err_rate_list).T
```

```

87 # return random-shuffled mini-batches
88 def random_mini_batches(image, label, mini_batch_size=256, one_batch=False):
89     m = image.shape[0] # number of training examples
90     mini_batches = []
91
92     # Shuffle (image, label)
93     permutation = list(np.random.permutation(m))
94     shuffled_image = image[permutation, :, :, :]
95     shuffled_label = label[permutation]
96
97     # extract only one batch
98     if one_batch:
99         mini_batch_image = shuffled_image[0: mini_batch_size, :, :, :]
100         mini_batch_label = shuffled_label[0: mini_batch_size]
101         return mini_batch_image, mini_batch_label
102
103     # Partition (shuffled_image, shuffled_Y). Minus the end case.
104     num_complete_minibatches = math.floor(m / mini_batch_size)
105     for k in range(0, num_complete_minibatches):
106         mini_batch_image = shuffled_image[k * mini_batch_size: k * mini_batch_size + mini_batch_size, :, :, :]
107         mini_batch_label = shuffled_label[k * mini_batch_size: k * mini_batch_size + mini_batch_size]
108         mini_batch = (mini_batch_image, mini_batch_label)
109         mini_batches.append(mini_batch)
110     # Handling the end case (last mini-batch < mini_batch_size)
111     if m % mini_batch_size != 0:
112         mini_batch_image = shuffled_image[num_complete_minibatches * mini_batch_size: m, :, :, :]
113         mini_batch_label = shuffled_label[num_complete_minibatches * mini_batch_size: m]
114         mini_batch = (mini_batch_image, mini_batch_label)
115         mini_batches.append(mini_batch)
116
117     return mini_batches

```

实验结果

加载数据集并进行处理

```

Loading MNIST data from files...
Load images from D:/wtx/machine-learning/ex1/mnist_data/train-images-idx3-ubyte, number: 60000, data shape: (60000, 28, 28)
Load images from D:/wtx/machine-learning/ex1/mnist_data/train-labels-idx1-ubyte, number: 60000, data shape: (60000,)
Load images from D:/wtx/machine-learning/ex1/mnist_data/t10k-images-idx3-ubyte, number: 10000, data shape: (10000, 28, 28)
Load images from D:/wtx/machine-learning/ex1/mnist_data/t10k-labels-idx1-ubyte, number: 10000, data shape: (10000,)
Got data...

Normalization and zero-padding...

The shape of training image with padding: (60000, 32, 32, 1)
The shape of testing image with padding: (10000, 32, 32, 1)
Finish data processing...

```


训练模型

```
Start training...
----- epoch1 begin -----
global learning rate: 0.05
learning rates in trainable layers: [1.31954867e-05 1.22353197e-05 1.88510729e-05 1.94846469e-05]
batch size: 256
Done, cost of epoch 1 : 670873.0466392617
0/1 error of training set: 3915 / 60000617
Time used: 291.01013588905334 sec
----- epoch1 end -----

----- epoch2 begin -----
global learning rate: 0.05
learning rates in trainable layers: [2.40907961e-04 2.03305658e-05 1.65692959e-05 4.80155308e-06]
batch size: 256
Done, cost of epoch 2 : 220625.76046242146
0/1 error of training set: 1982 / 600002146
Time used: 313.1585421562195 sec
----- epoch2 end -----
```

测试模型

```
----- epoch6 begin -----
global learning rate: 0.01
learning rates in trainable layers: [5.70590825e-04 7.32549595e-06 4.68012735e-06 7.17054754e-07]
batch size: 256
Done, cost of epoch 6 : 119039.96486393375
0/1 error of training set: 1317 / 600003375
Time used: 289.01421427726746 sec
----- epoch6 end -----

Finished training, the total training time is 1804.8934412002563s

Start testing...
error rate: 0.022
Finished testing, the accuracy is 0.978
```

正确率为 97.8%，正确率较低的原因可能会训练模型时训练论述较低。但由于电脑配置较低，在轮数较高时会因为系统内存不足导致整个程序卡死，因此不得不降低轮数。（在只运行 vcode 的情况下，每一轮中，系统内存使用最高达到了 96%）